

# **30 Python Libraries** **to (Hugely) Boost** **Your Data Science** **Productivity**

**Avi Chawla**



[avichawla.substack.com](http://avichawla.substack.com)




**Data Science is much  
more than Pandas,  
NumPy and  
Sklearn.**

Here are **30 open-source  
libraries** to upgrade  
your data game.





# 1. YellowBrick



A suite of visualization and diagnostic tools for **faster model selection.**


Matplotlib

Sklearn





## 2. PyCaret

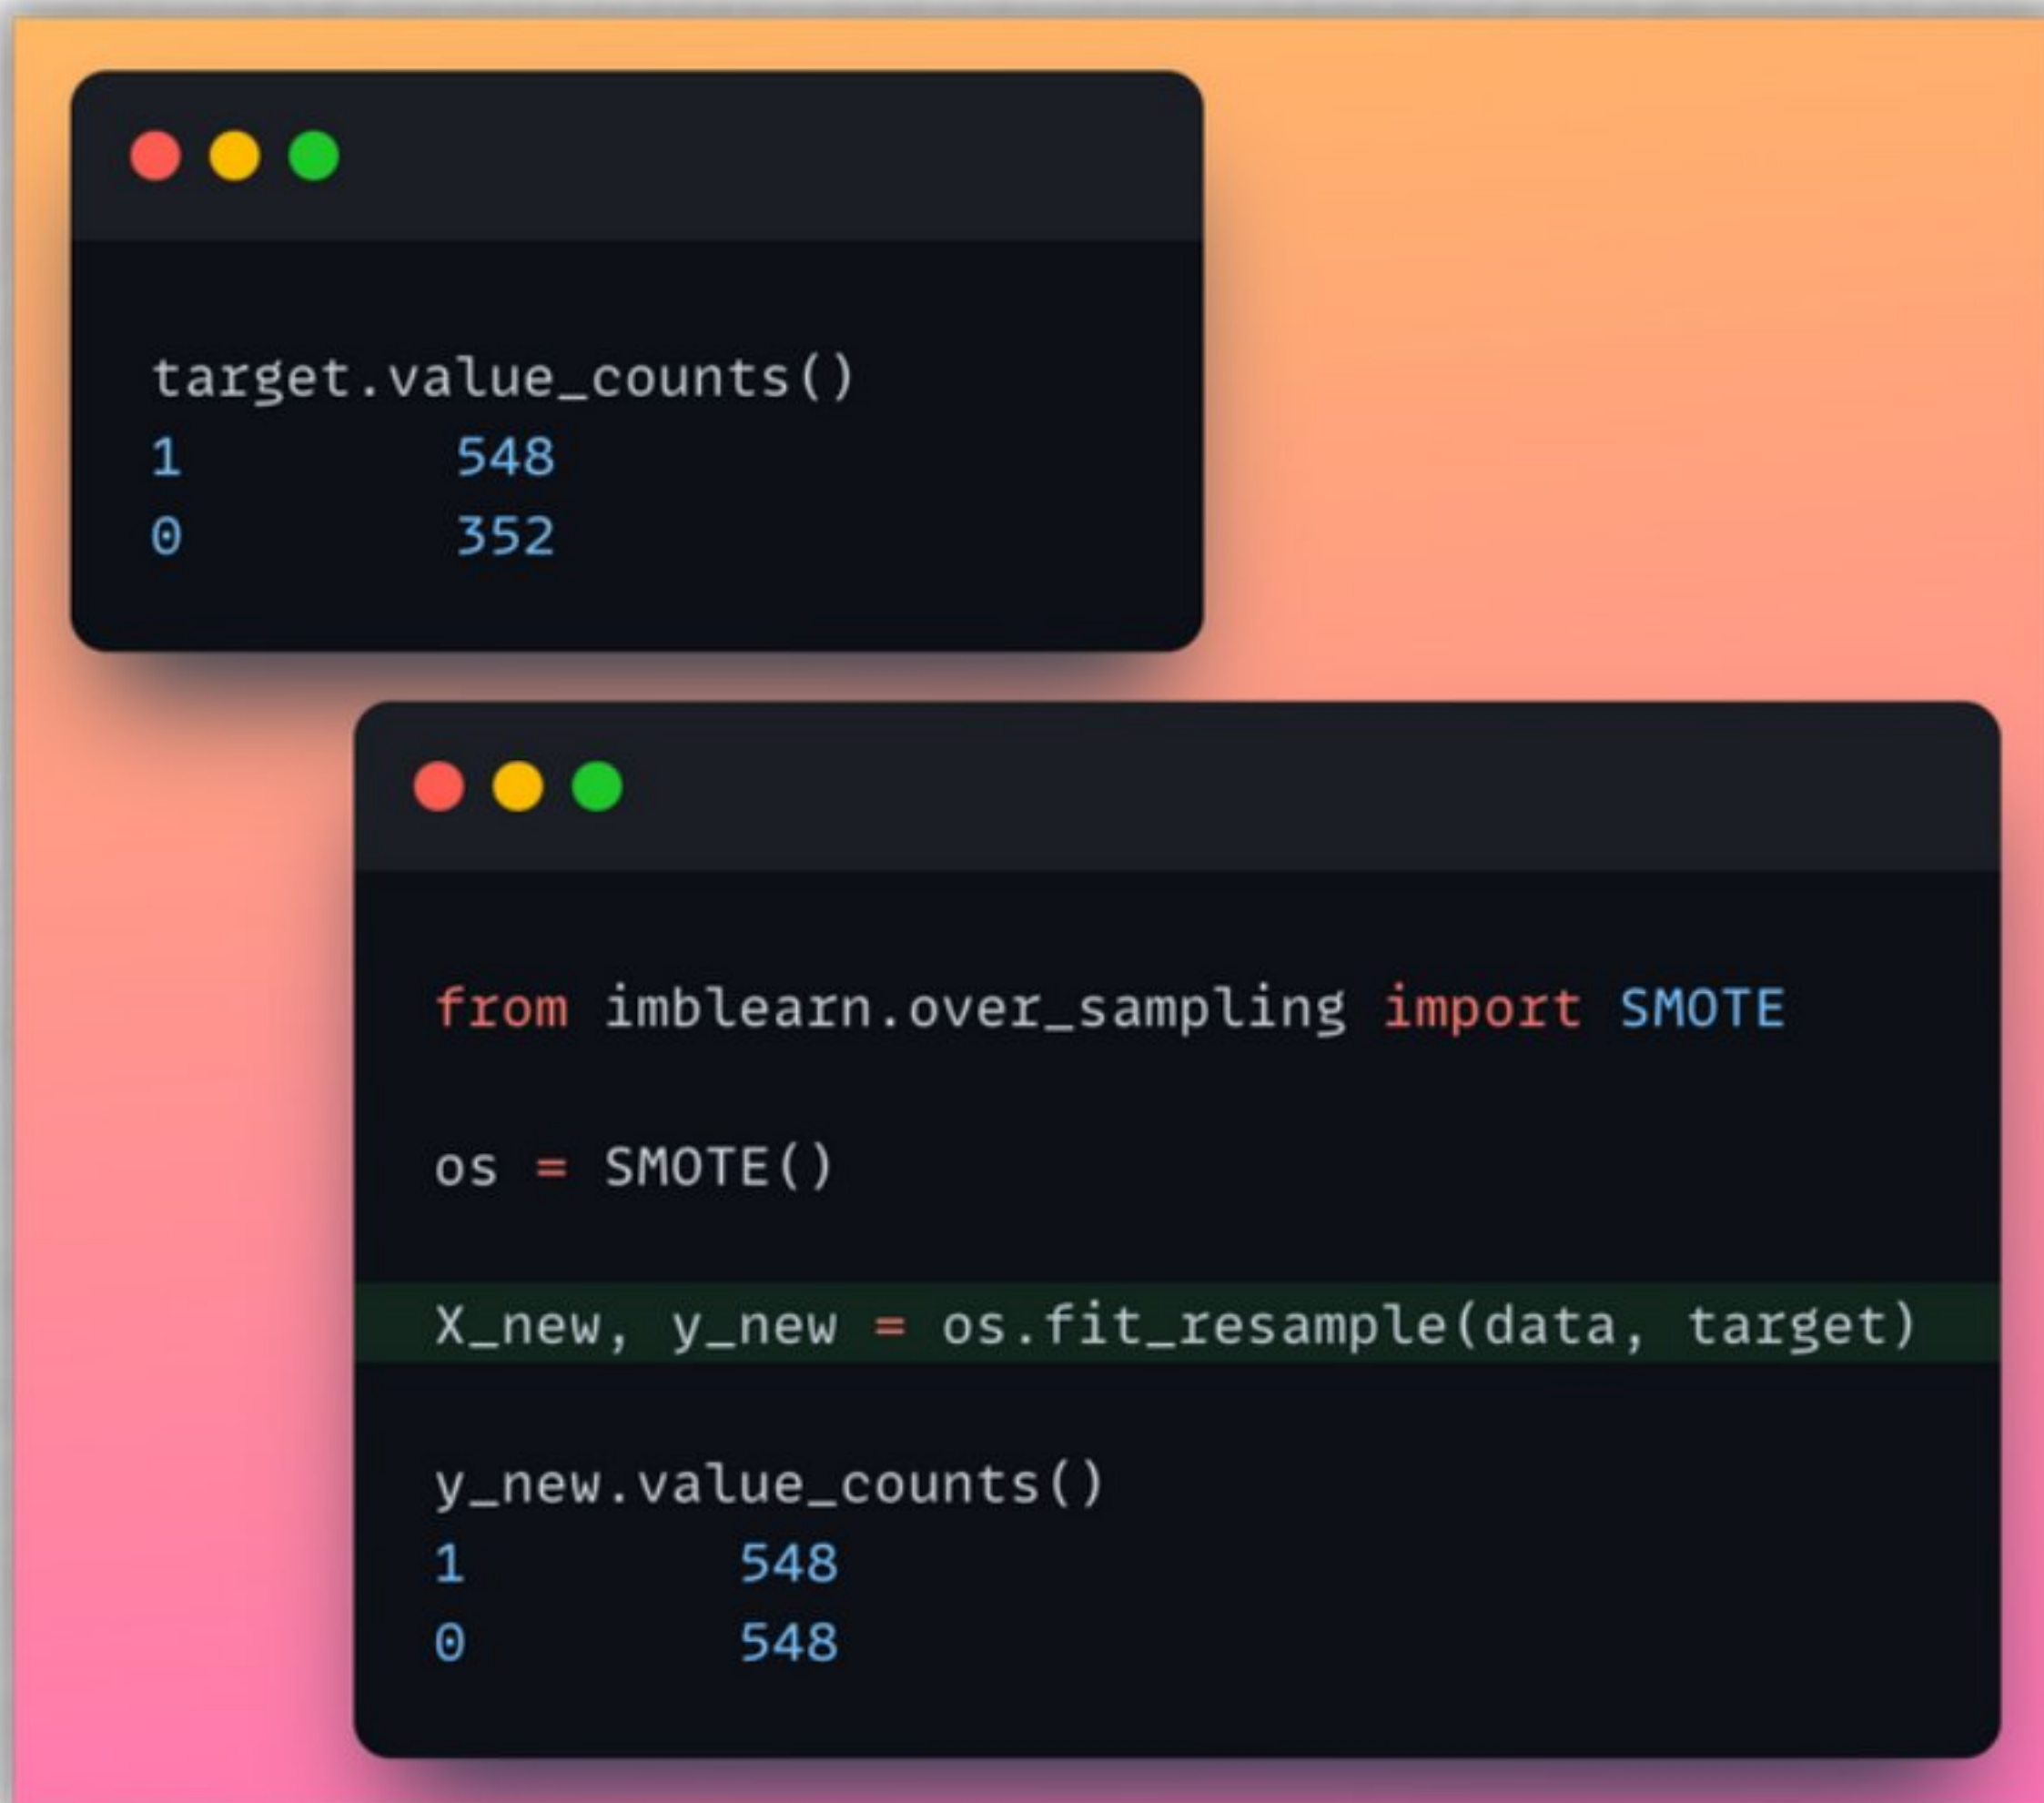


Automate ML workflows  
with this **low-code  
library**.





## 3. imbalanced-learn



```
target.value_counts()
1      548
0      352
```

```
from imblearn.over_sampling import SMOTE

os = SMOTE()

X_new, y_new = os.fit_resample(data, target)


y_new.value_counts()
1      548
0      548
```



A variety of methods  
to handle **class**  
**imbalance**.



## 4. Modin




Boost Pandas' performance **up to 70x** by modifying the import.





## 5. SHAP

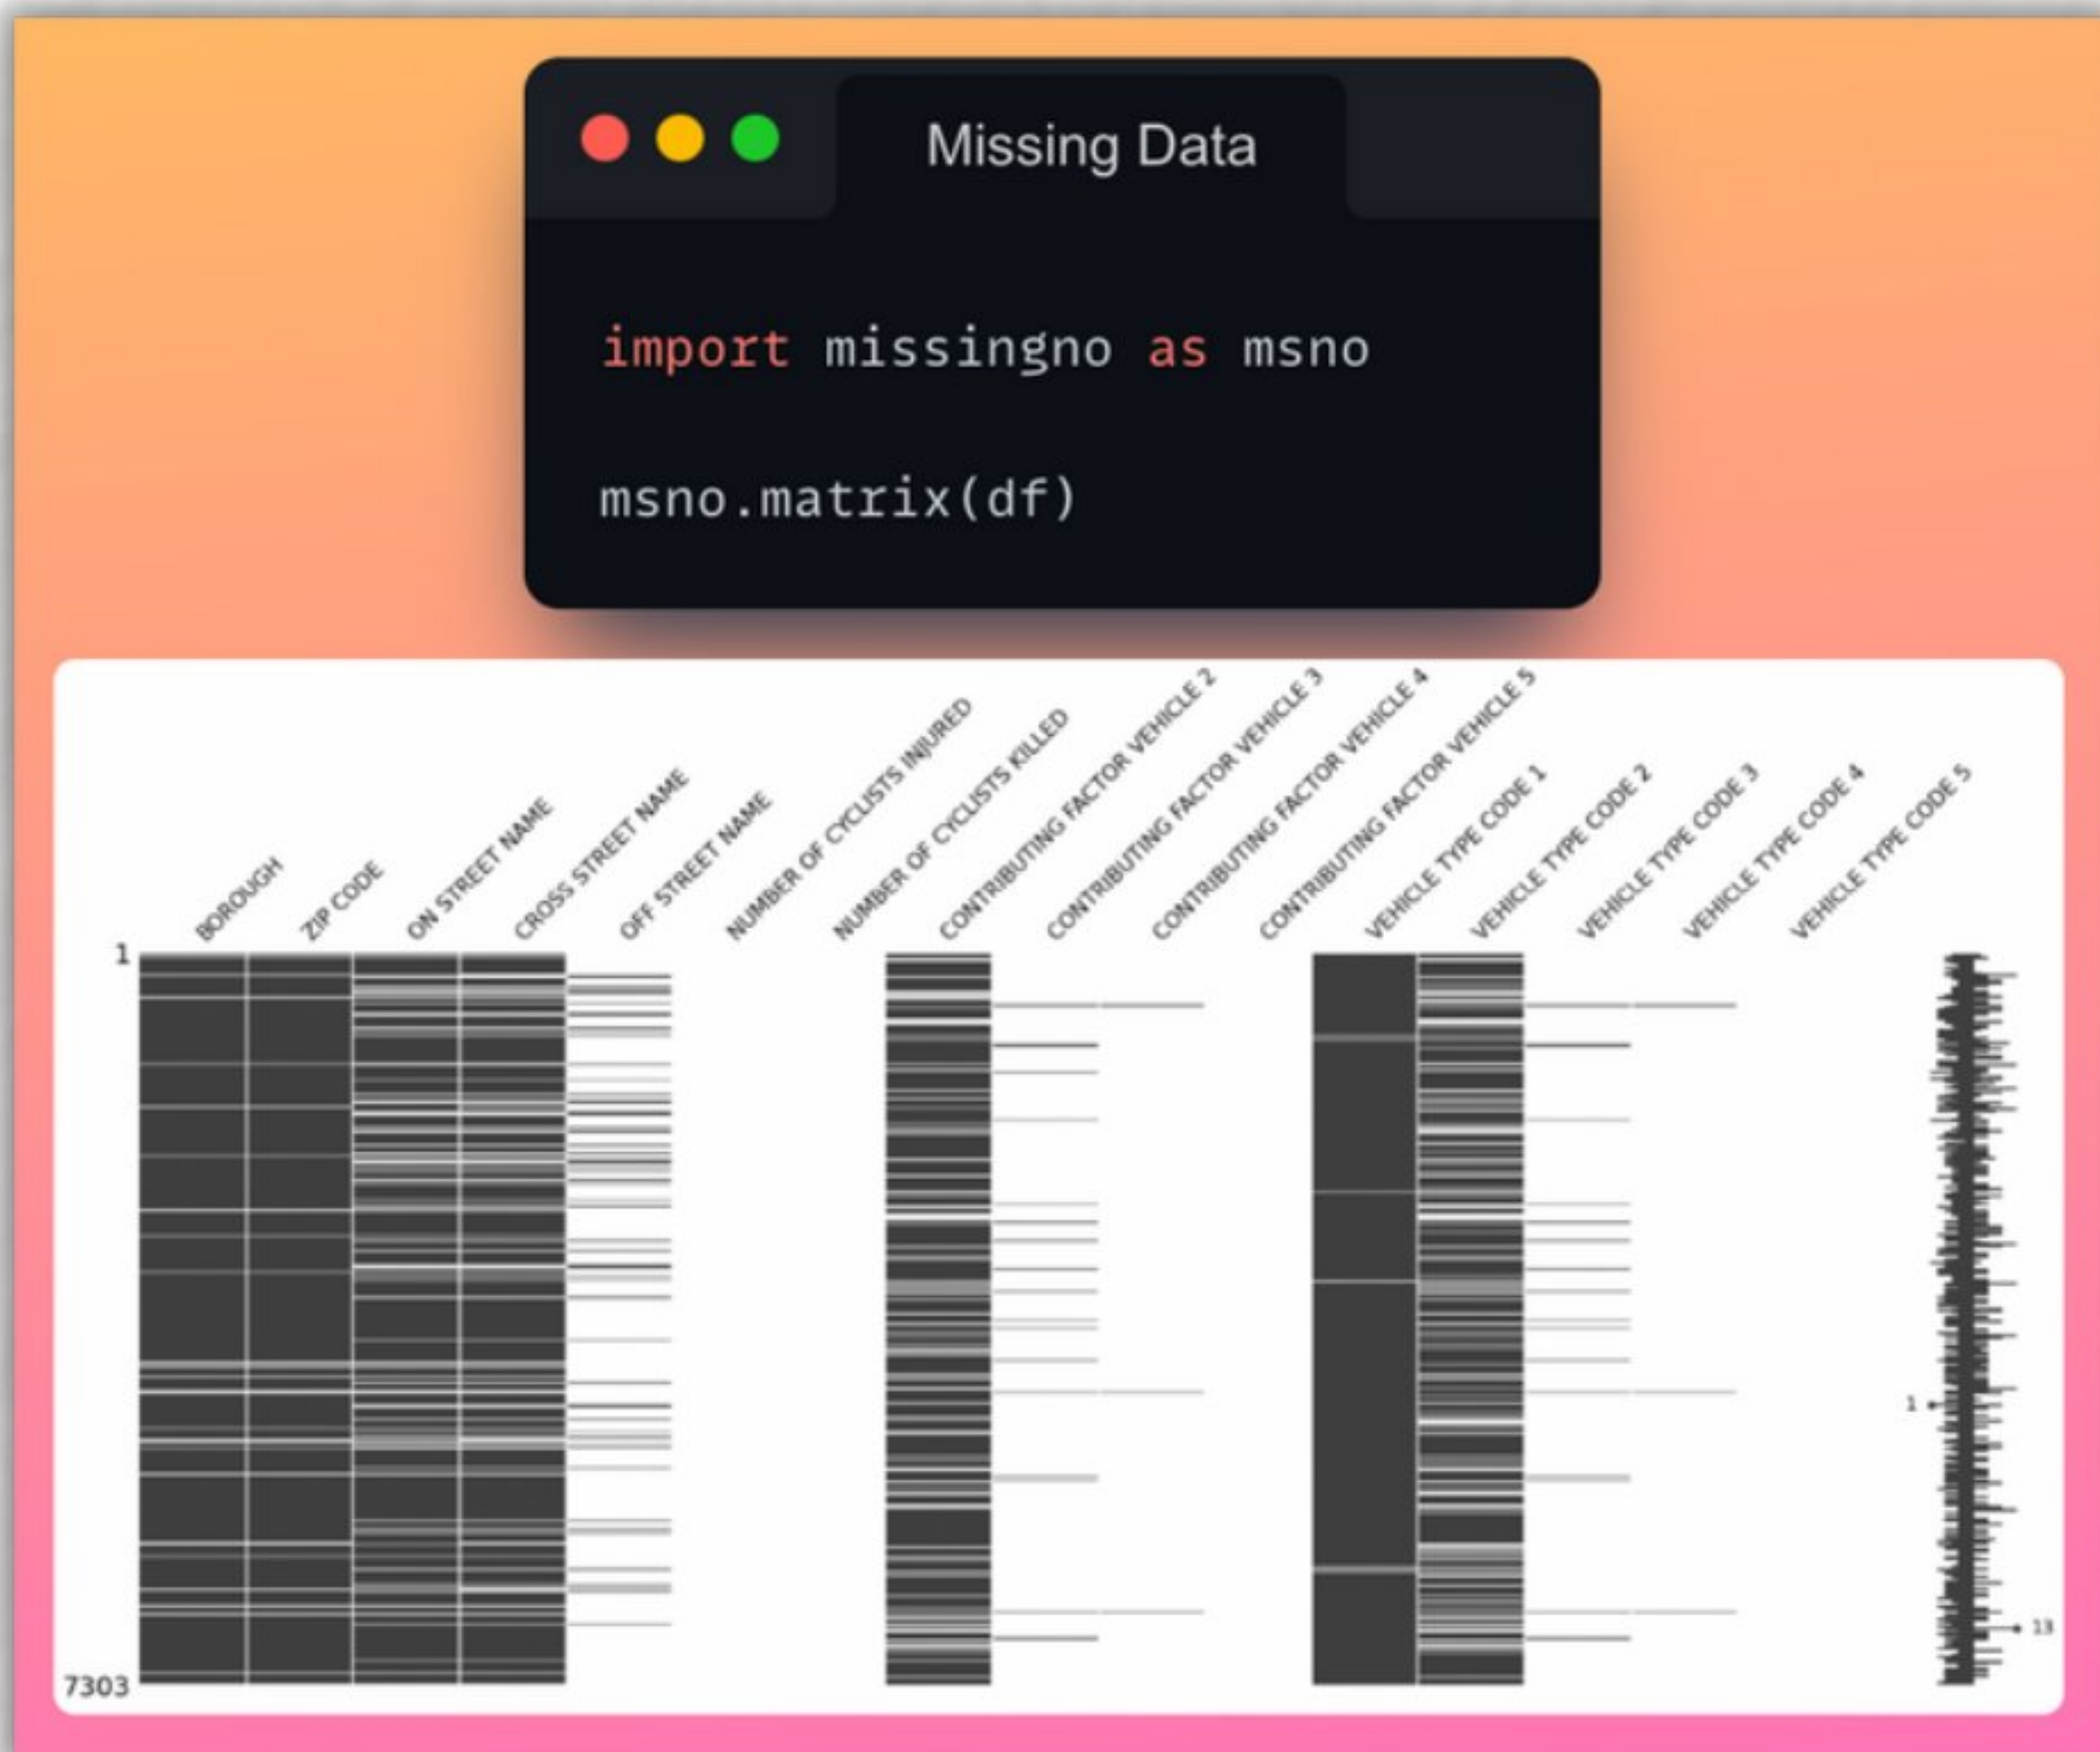


Explain the output of any ML model in few lines of code.






## 6. Missingno




Visualize missing values  
in your dataset  
with ease.





# 7. Prophet



Produce **high-quality forecasts** on time-series data.



## 8. Parallel-Pandas

```
from parallel_pandas import ParallelPandas  
  
ParallelPandas.initialize()
```

```
- df.apply()    #Pandas  
+ df.p_apply() #Parallel
```

```
- df.map()      #Pandas  
+ df.p_map()   #Parallel
```

```
- df.mean()     #Pandas  
+ df.p_mean()  #Parallel
```

```
- df.describe() #Pandas  
+ df.p_describe() #Parallel
```

 Parallelize Pandas across all CPU cores for **faster computation**.



## 9. Featuretools



 **Automated feature engineering for ML models.**



# 10. Lazy Predict

```
from lazypredict.Supervised import LazyRegressor

reg = LazyRegressor()
reg.fit(X_train, X_test, y_train, y_test)
```

```
| Model | R-Squared | RMSE | Time Taken |
|:-----:|:-----:|:-----:|:-----:|
| SVR | 0.877199 | 2.62054 | 0.0330021 |
| KNeighborsRegressor | 0.826307 | 3.1166 | 0.0179954 |
| MLPRegressor | 0.750536 | 3.73503 | 0.725997 |
| ... | ... | ... | ... |
| LinearRegression | 0.71753 | 3.97444 | 0.0190051 |
| DummyRegressor | -0.02157 | 7.55832 | 0.0140116 |
```

 Train **30 machine learning models** in one line of code.




# 11. mlxtend

```
from mlxtend.plotting import plot_decision_regions

model = LogisticRegression().fit(X, y)

plot_decision_regions(X, y, model)
```



A collection of utility functions for **processing, evaluating, visualizing** models.



# 12. Vaex

```
%%time  
df
```

CPU times: user 8  $\mu$ s, sys: 0 ns, total: 8  $\mu$ s  
Wall time: 16.7  $\mu$ s


| #             | vendor_id | pickup_datetime                  | dropoff_datetime                 | passenger_count | payment_type | trip_distance      | pickup_longitude   | pickup_latitude    |
|---------------|-----------|----------------------------------|----------------------------------|-----------------|--------------|--------------------|--------------------|--------------------|
| 0             | VTS       | 2009-01-04<br>02:52:00.000000000 | 2009-01-04<br>03:02:00.000000000 | 1               | CASH         | 2.630000114440918  | -73.99195861816406 | 40.72156524658203  |
| 1             | VTS       | 2009-01-04<br>03:31:00.000000000 | 2009-01-04<br>03:38:00.000000000 | 3               | Credit       | 4.550000190734863  | -73.98210144042969 | 40.736289978027344 |
| 2             | VTS       | 2009-01-03<br>15:43:00.000000000 | 2009-01-03<br>15:57:00.000000000 | 5               | Credit       | 10.350000381469727 | -74.0025863647461  | 40.73974609375     |
| 3             | DDS       | 2009-01-01<br>20:52:58.000000000 | 2009-01-01<br>21:14:00.000000000 | 1               | CREDIT       | 5.0                | -73.9742660522461  | 40.79095458984375  |
| 4             | DDS       | 2009-01-24<br>16:18:23.000000000 | 2009-01-24<br>16:24:56.000000000 | 1               | CASH         | 0.4000000059604645 | -74.00157928466797 | 40.719383239746094 |
| ...           | ...       | ...                              | ...                              | ...             | ...          | ...                | ...                | ...                |
| 1,173,057,922 | VTS       | 2015-12-31<br>23:59:56.000000000 | 2016-01-01<br>00:08:18.000000000 | 5               | 1            | 1.2000000476837158 | -73.99381256103516 | 40.72087097167969  |
| 1,173,057,923 | CMT       | 2015-12-31<br>23:59:58.000000000 | 2016-01-01<br>00:05:19.000000000 | 2               | 2            | 2.0                | -73.96527099609375 | 40.76028060913086  |
| 1,173,057,924 | CMT       | 2015-12-31<br>23:59:59.000000000 | 2016-01-01<br>00:12:55.000000000 | 2               | 2            | 3.799999952316284  | -73.98729705810547 | 40.739078521728516 |
| 1,173,057,925 | VTS       | 2015-12-31<br>23:59:59.000000000 | 2016-01-01<br>00:10:26.000000000 | 1               | 2            | 1.9600000381469727 | -73.99755859375    | 40.72569274902344  |
| 1,173,057,926 | VTS       | 2015-12-31<br>23:59:59.000000000 | 2016-01-01<br>00:21:30.000000000 | 1               | 1            | 1.059999942779541  | -73.9843978881836  | 40.76725769042969  |



High performance package  
for **lazy Out-of-Core  
DataFrames.**




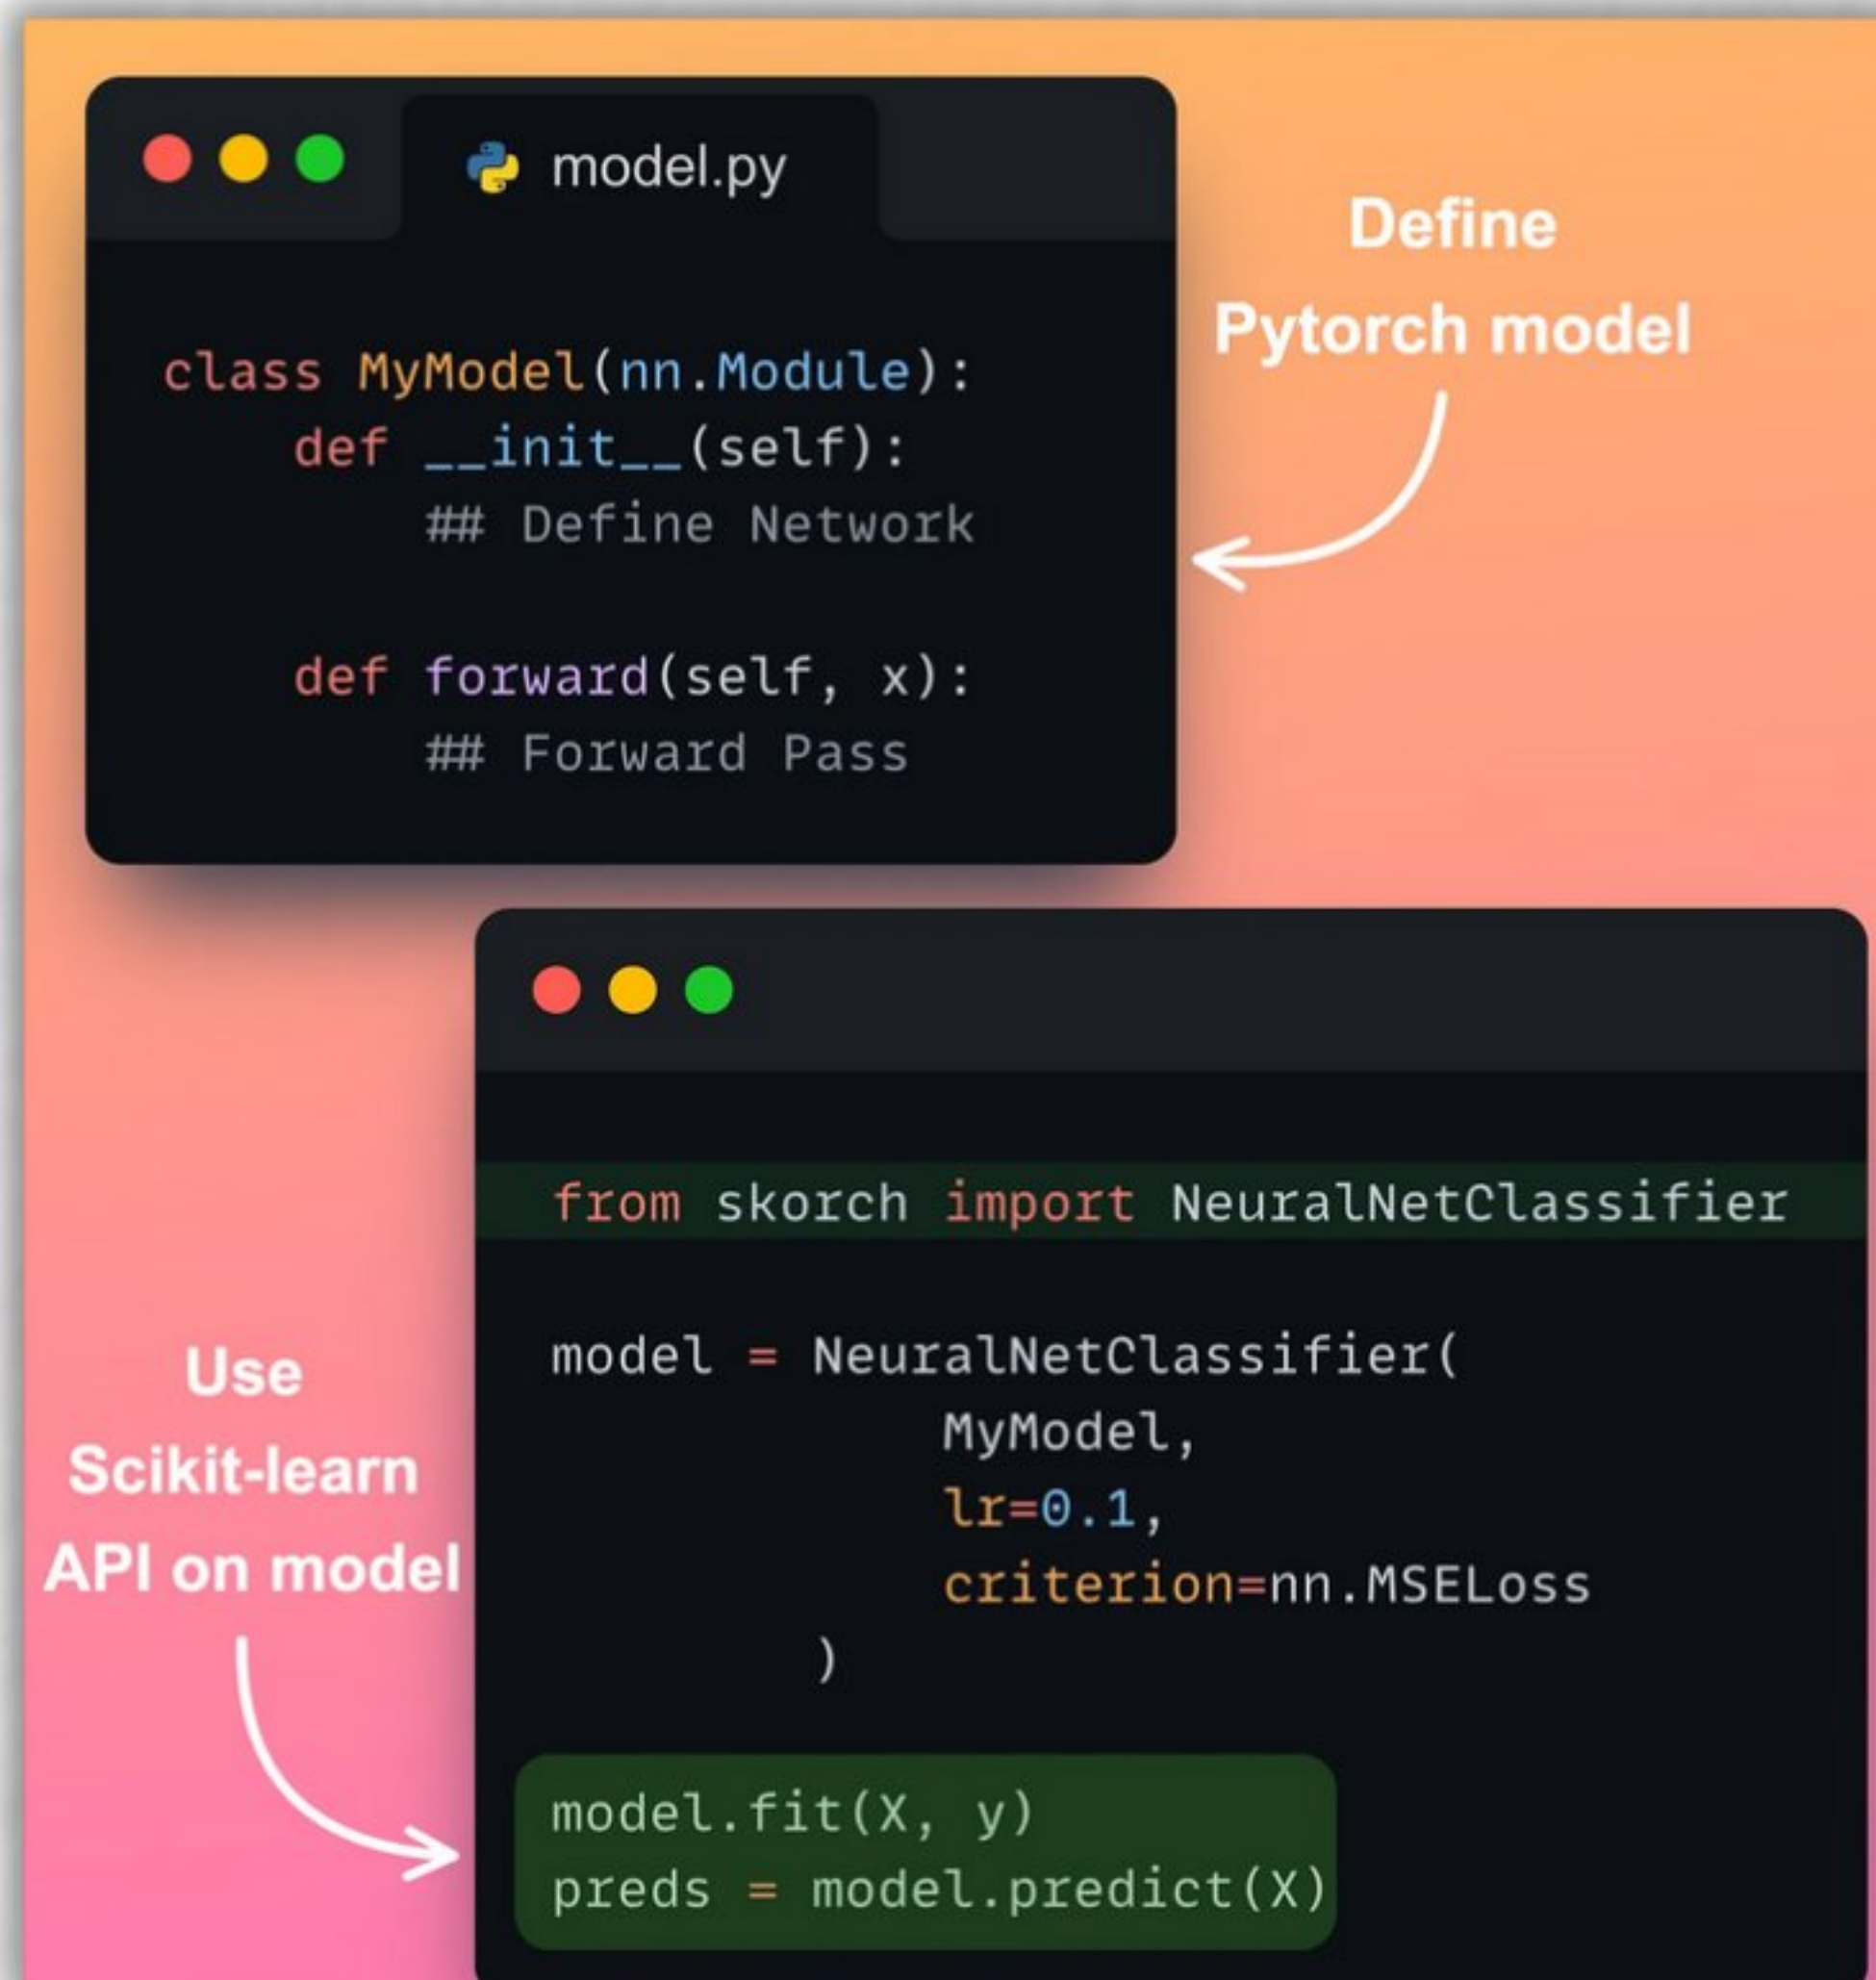
# 13. SweetViz



In-depth EDA report  
in two lines  
of code.




# 14. Skorch





# 15. Faiss



The image shows two code snippets side-by-side, each in a dark-themed editor window. The top window, titled 'sklearn.py', shows a KMeans model training in 162 seconds. The bottom window, titled 'faiss.py', shows the same KMeans model training in 7.8 seconds using Faiss. A pink arrow points from the text '~20x Faster' to the faiss training time.

```
sklearn.py
```

```
from sklearn.cluster import KMeans

kmeans = KMeans(8).fit(x_train)
# Training Time: 162s
```

**~20x Faster**

```
faiss.py
```

```
import faiss

kmeans = faiss.Kmeans(d=1024, k=8)
kmeans.train(x_train)
# Training Time: 7.8s
```

 Efficient algorithms for **similarity search** and **clustering** dense vectors.





# 16. statsmodel

```
import statsmodels.api as sm

results = sm.OLS(y, X).fit()
print(results.summary())
```

OLS Regression Results

|                   |                  |                     |          |
|-------------------|------------------|---------------------|----------|
| Dep. Variable:    | y                | R-squared:          | 0.178    |
| Model:            | OLS              | Adj. R-squared:     | 0.161    |
| Method:           | Least Squares    | F-statistic:        | 10.51    |
| Date:             | Wed, 02 Nov 2022 | Prob (F-statistic): | 7.41e-05 |
| Time:             | 17:12:45         | Log-Likelihood:     | -20.926  |
| No. Observations: | 100              | AIC:                | 47.85    |
| Df Residuals:     | 97               | BIC:                | 55.67    |
| Df Model:         | 2                |                     |          |
| Covariance Type:  | nonrobust        |                     |          |

|       | coef   | std err | t      | P> t  | [0.025 | 0.975] |
|-------|--------|---------|--------|-------|--------|--------|
| const | 1.4713 | 0.075   | 19.579 | 0.000 | 1.322  | 1.620  |
| x1    | 0.1045 | 0.105   | 1.000  | 0.320 | -0.103 | 0.312  |
| x2    | 0.4831 | 0.107   | 4.503  | 0.000 | 0.270  | 0.696  |

|                |        |                   |        |
|----------------|--------|-------------------|--------|
| Omnibus:       | 39.684 | Durbin-Watson:    | 1.848  |
| Prob(Omnibus): | 0.000  | Jarque-Bera (JB): | 6.503  |
| Skew:          | 0.096  | Prob(JB):         | 0.0387 |
| Kurtosis:      | 1.766  | Cond. No.         | 5.09   |

 **Statistical testing and data exploration at fingertips.**






# 17. Pandas-Profiling

```
from pandas_profiling import ProfileReport

profile = ProfileReport(iris_data,
                        title="Pandas Profiling Report")


profile.to_widgets()
```



Generate a high-level **EDA report** of your data in no time.



# 18. Streamlit



→ Create and host data-based  
Python web apps  
in few lines  
of code.





# 19. Category-encoders

```
import category_encoders as ce

ce_be = ce.BinaryEncoder(cols=['class'])

ce_be.fit_transform(data["class"])
```

|   | class_0 | class_1 | class_2 |
|---|---------|---------|---------|
| 0 | 0       | 0       | 1       |
| 1 | 0       | 1       | 0       |
| 2 | 0       | 1       | 1       |
| 3 | 1       | 0       | 0       |
| 4 | 0       | 0       | 1       |



**Binary**


|   | gender | class |
|---|--------|-------|
| 0 | Male   | A     |
| 1 | Female | B     |
| 2 | Male   | C     |
| 3 | Female | D     |
| 4 | Female | A     |



**Over 15** categorical  
data encoders.



## 20. DuckDB



Run **SQL queries**  
on **DataFrame**.



## 21. PandasML

```
import pandas_ml as pdml

df = pdml.ModelFrame(X, y)

train, test = df.model_selection.
               train_test_split()

reg = df.linear_model.LinearRegression()

train.fit(reg)

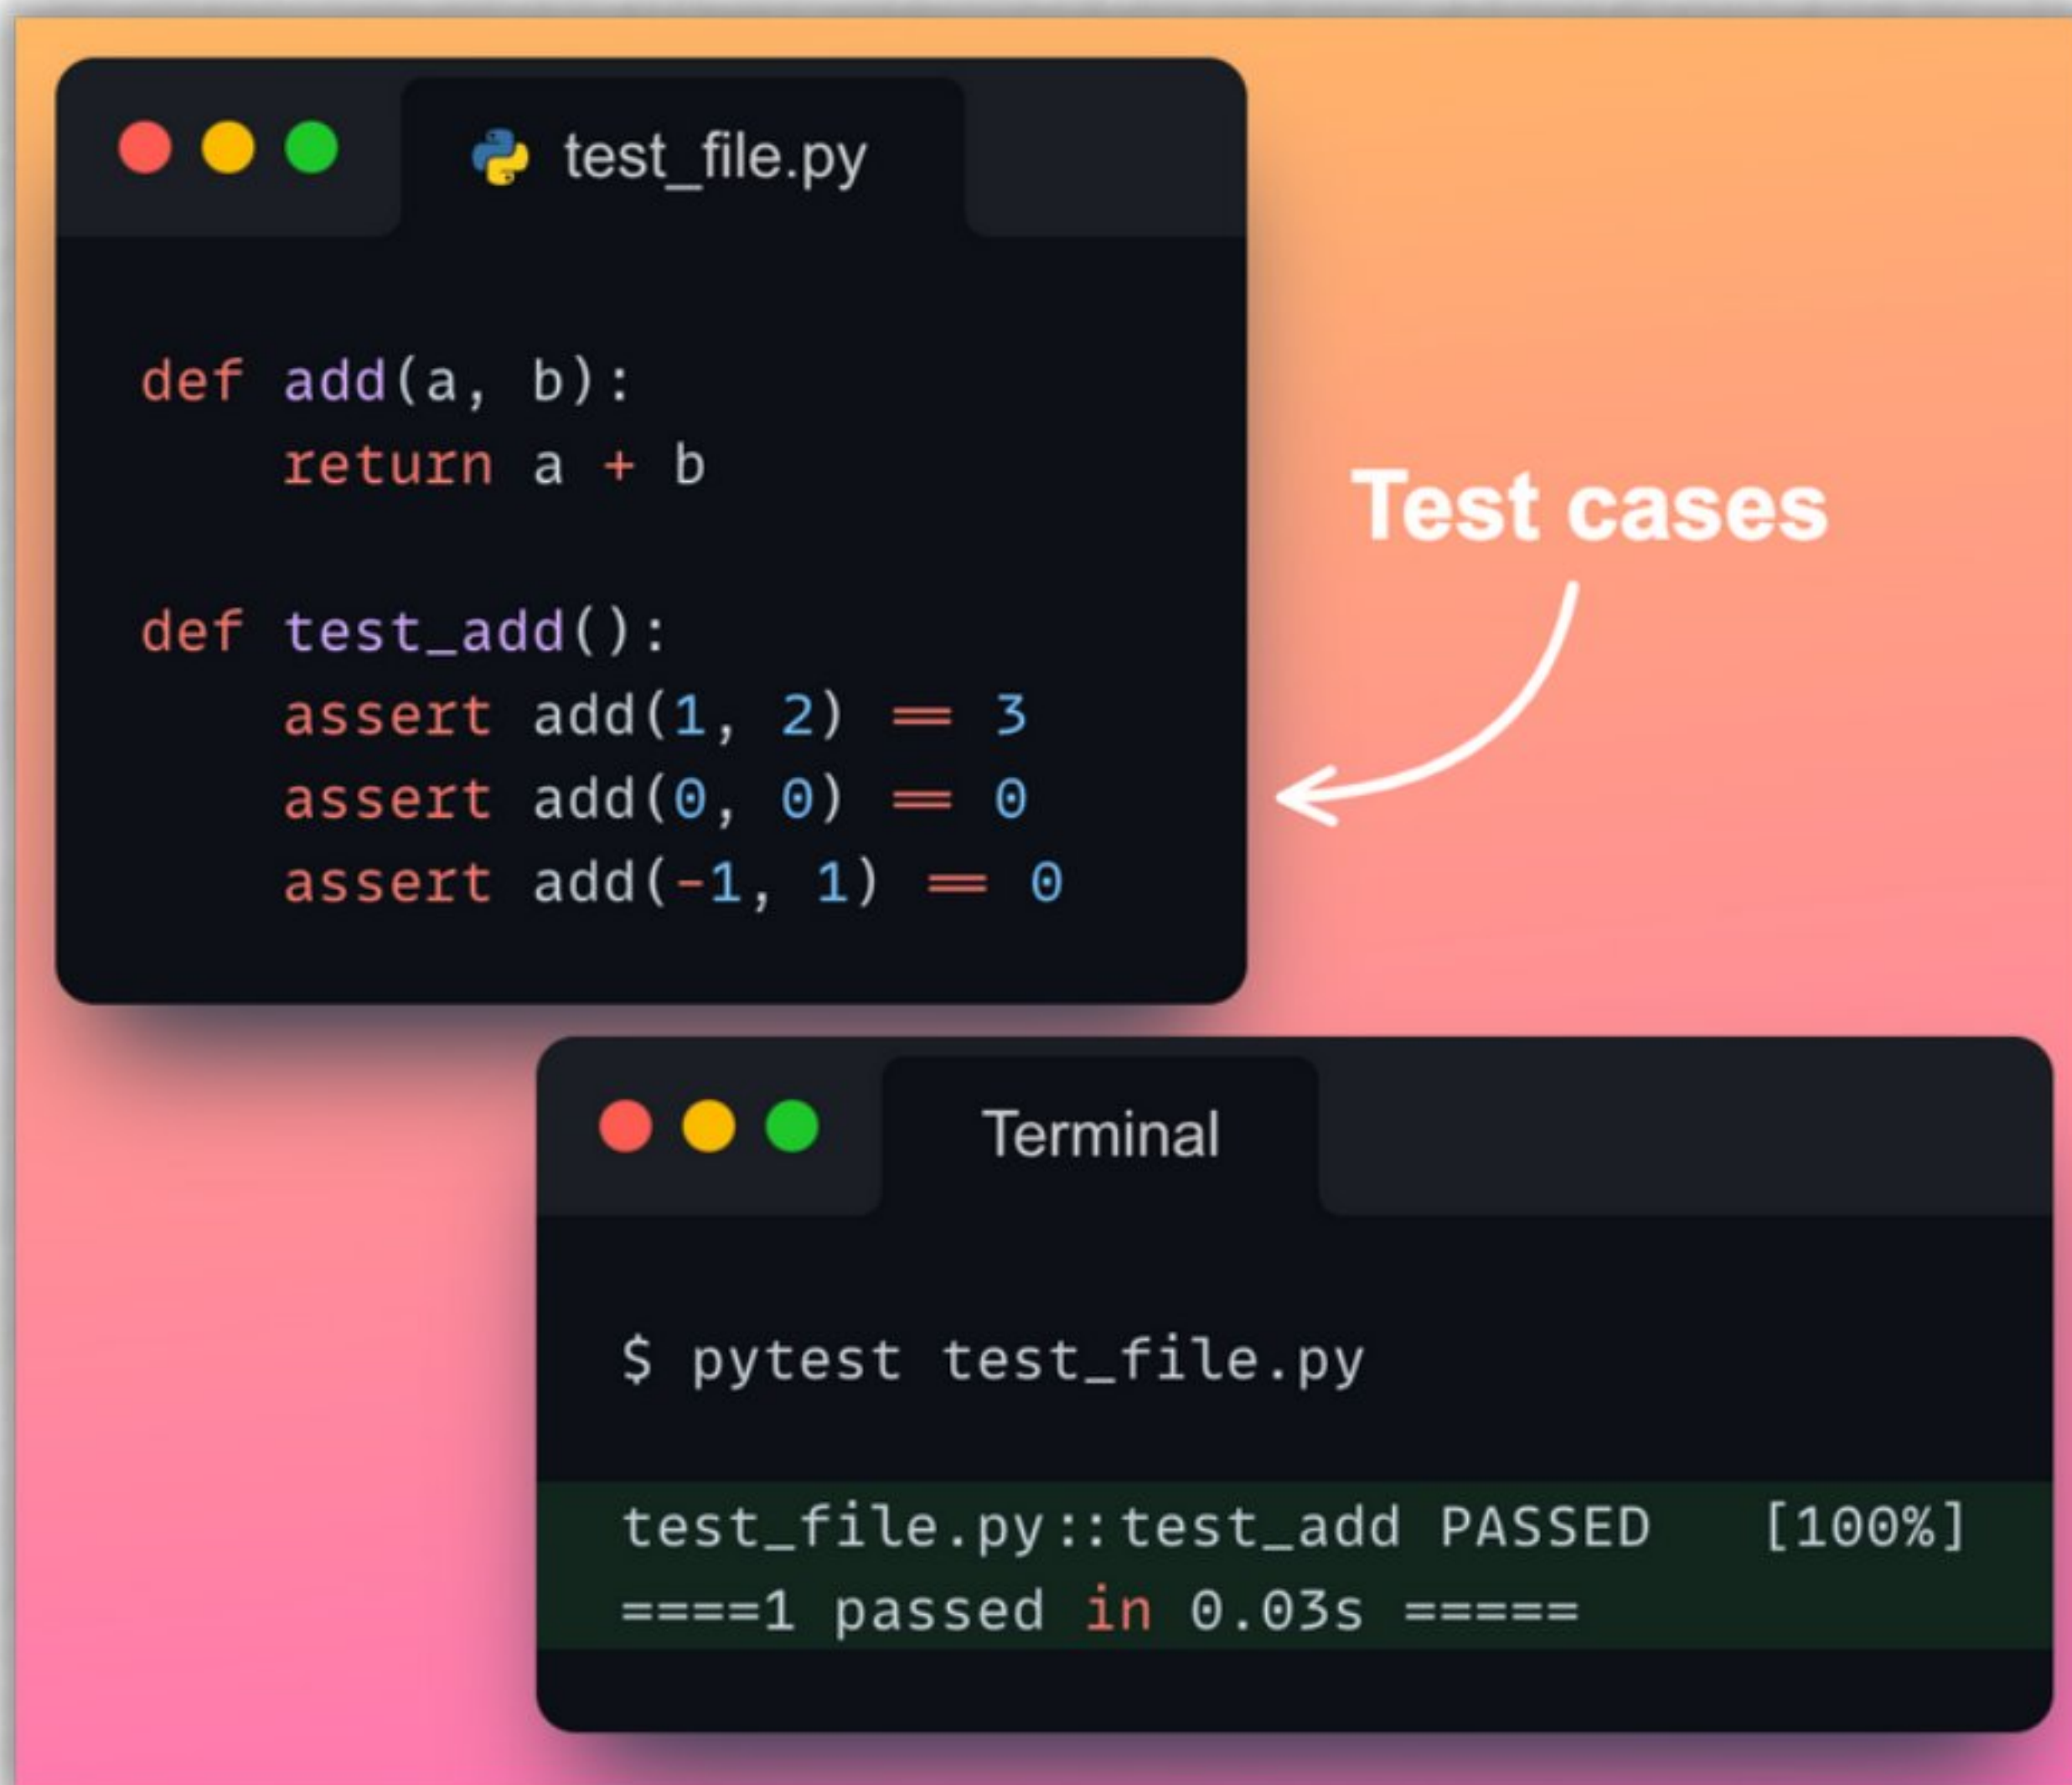
test.predict(reg)
```

→ **Pandas** data wrangling +  
**Sklearn** algorithms +  
**Matplotlib** visualization.





## 22. Pytest



 An elegant testing framework to test your code.





## 23. Numexpr

```
import numpy as np
import numexpr as ne
```

```
a = np.random.random(10**7)
b = np.random.random(10**7)
```

```
%timeit np.cos(a) + np.sin(b)
```

142 ms ± 257 µs per loop

```
%timeit ne.evaluate("cos(a) + sin(b)")
```

32.5 ms ± 229 µs per loop **~5x Faster**

 Parallelize NumPy to  
all CPU cores for  
**20x speedup.**





## 24. CSV-Kit

**data.csv**

| Name  | Marks | Grade |
|-------|-------|-------|
| Joe   | 95    | A     |
| Hanna | 89    | B     |
| Chris | 92    | A     |
| Julie | 94    | A     |

**Column Names**

```
$ csvcut -n data.csv
1: Name
2: Marks
3: Grade
```

**Column Stats**

```
$ csvstat data.csv
2. "Marks"
Type of data:      Number
Contains null values: False
Unique values:     4
Smallest value:    89
Largest value:     95
Sum:               370
Mean:              92.5
Median:            93
StDev:             2.646
```

**Query**

```
$ csvsql --query "select * from data where Marks>90" data.csv
```


| Name  | Marks | Grade |
|-------|-------|-------|
| Joe   | 95    | A     |
| Chris | 92    | A     |
| Julie | 94    | A     |

Explore, query and describe CSV files from terminal.





## 25. PivotTableJS



The screenshot shows a Jupyter Notebook interface with a dark-themed code editor and a light-themed output area. The code editor displays the following Python code:

```
from pivottablejs import pivot_ui
pivot_ui(df)
```

The output area shows a pivot table with the following configuration:

- Table: Employee
- Count: Employee\_City
- Employment\_Status: Employment\_Status

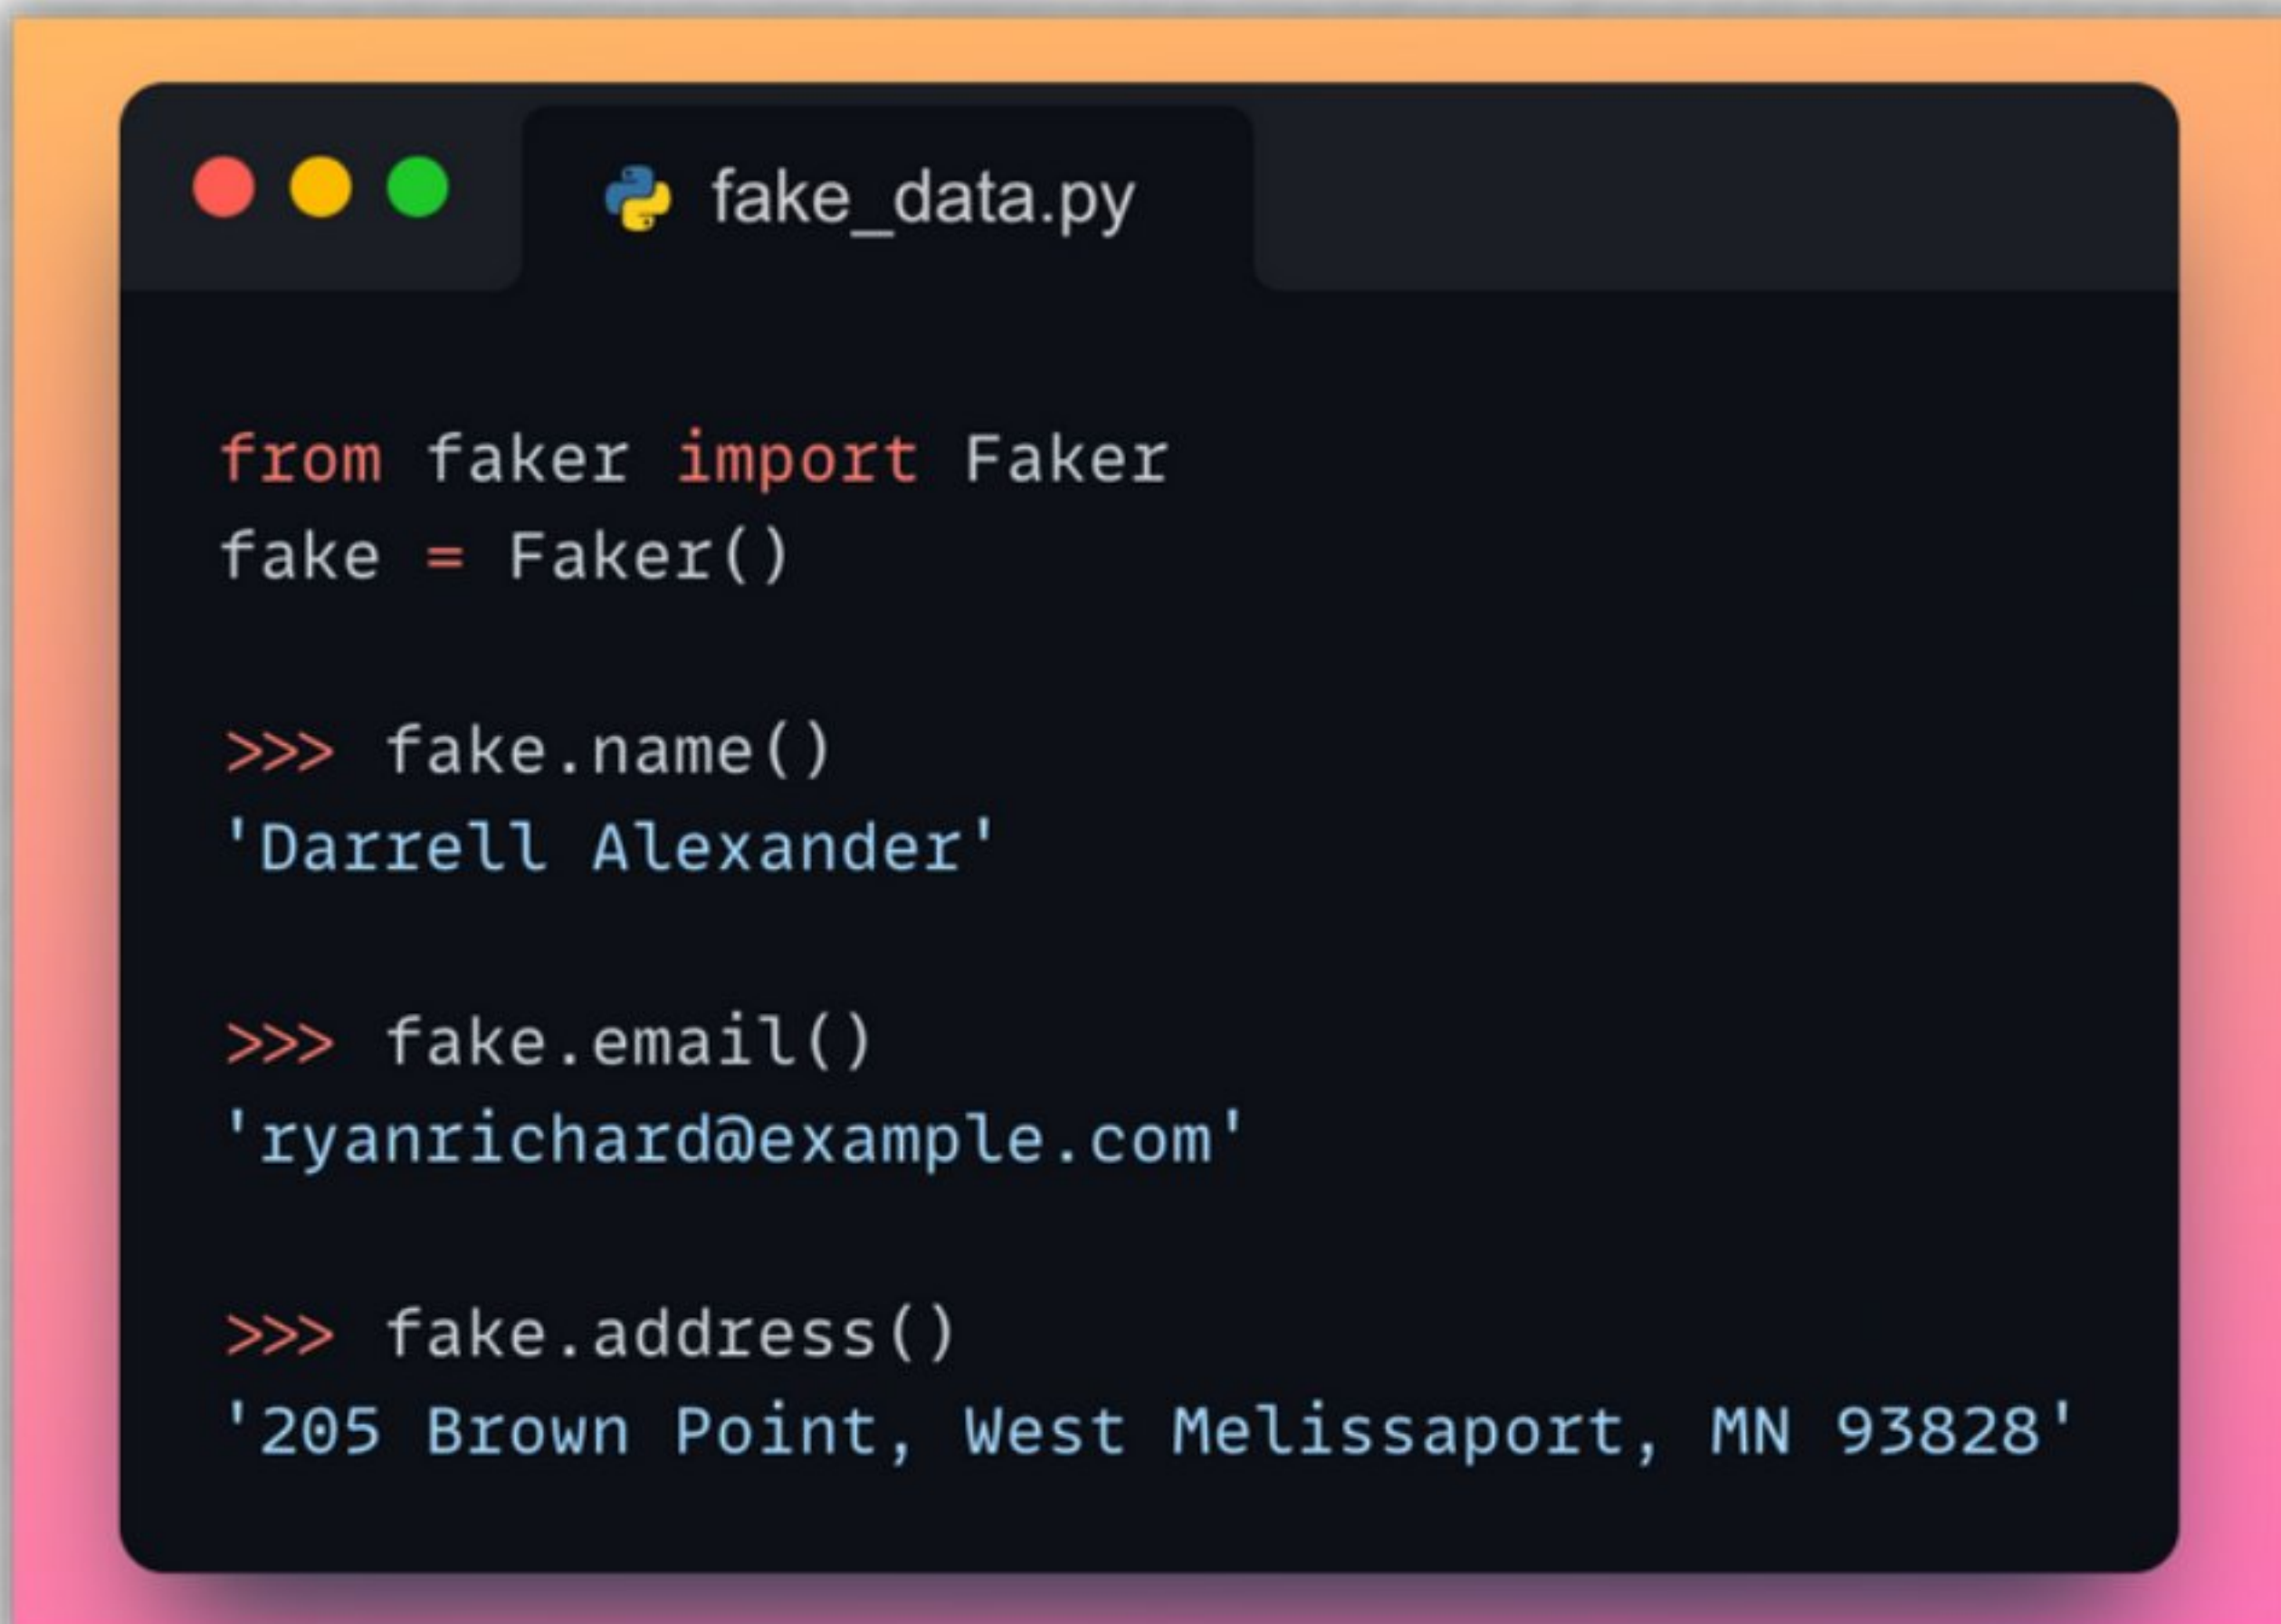
The pivot table displays the following data:

| Employee_City     | Employment_Status | Full Time | Intern | Totals |
|-------------------|-------------------|-----------|--------|--------|
| Aliciafort        |                   | 89        | 24     | 113    |
| Kristaburgh       |                   | 77        | 20     | 97     |
| New Cindychester  |                   | 82        | 24     | 106    |
| New Russellton    |                   | 73        | 20     | 93     |
| North Melissafurt |                   | 62        | 16     | 78     |
| Ricardomouth      |                   | 81        | 25     | 106    |
| Wardfort          |                   | 82        | 14     | 96     |
| West Jamesview    |                   | 94        | 26     | 120    |
| Whitakerbury      |                   | 66        | 21     | 87     |
| Whiteside         |                   | 85        | 19     | 104    |
| Totals            |                   | 791       | 209    | 1,000  |

→ **Drag-n-drop** tools to group, pivot, plot dataframe.



## 26. Faker



```
from faker import Faker
fake = Faker()

>>> fake.name()
'Darrell Alexander'

>>> fake.email()
'ryanrichard@example.com'

>>> fake.address()
'205 Brown Point, West Melissaport, MN 93828'
```

 Generate **fake yet meaningful data** in seconds.



## 27. Icecream

```
Print

def func(arr, n):
    print("arr =", arr, "n =", n)

func([1,2,3], 2)
## arr = [1, 2, 3] n = 2
```

```
Icecream

from icecream import ic

def func(arr, n):
    ic(arr, n)

func([1,2,3], 2)
## ic| arr: [1, 2, 3], n: 2
```

 **Don't debug with `print()`. Use icecream.**



## 28. Pyforest

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import sys

from sklearn.linear_model import LinearRegression
```




```
from pyforest import *

pd.read_csv("file.csv") ✓
np.array([1,2,3]) ✓
sys.path ✓
LinearRegression() ✓
```

 No need to write imports.  
**Automatic package  
import.**



## 29. PySnooper



The image shows a code editor window titled 'py-snooper.py' and a terminal window below it. The code editor contains the following Python code:

```
1 import pysnooper
2
3 @pysnooper.snoop()
4 def add_sub(a, b):
5
6     add = a+b
7     sub = a-b
8
9     return (add, sub)
10
11 add_sub(9, 5)
```

An arrow labeled 'Add Decorater' points to line 3. The terminal window shows the output of running the code:

```
$ python py-snooper.py

Starting var:.. a = 9
Starting var:.. b = 5
call
line
New var:..... add = 14
line
New var:..... sub = 4
line
Return value:.. (14, 4)

4 def add_sub(a, b):
6     add = a+b
7     sub = a-b
9     return (add, sub)
```

An arrow labeled 'Debugging Output' points to the terminal output.

Profile your code. Track new variables, and their updates.



## 30. Sidetable




Supercharge Pandas' **value\_counts()** method.



# Hope that helped.



Checkout my daily newsletter to learn something new about **Python** and **Data Science** everyday.

 [avichawla.substack.com](https://avichawla.substack.com)

Connect with me on **LinkedIn**.

<https://www.linkedin.com/in/avi-chawla>

